

# Assuring business value



Lucas Jellema, CTO of AMIS  
Fall 2018

## Monitoring

with Prometheus  
and Grafana

# Lucas Jellema

## Architect / Developer

1994 started in IT at Oracle

2002 joined AMIS

Currently CTO & Solution Architect



# Overview



- Observe [non-functional] behavior of business functions/applications in pnear] real-time
  - Availability and health
  - Performance
  - Access – as it should be for whom it should be
- Teams up with:
  - Profiling & Debugging – maximum context for activity spikes
  - Tracing – track paths through application [and platform & infra] stack
  - Logging – per application or platform component output for off-line processing
    - although log agents such as Elastic Stack Beats and Elasticsearch

# What is required for monitoring?



- Gather metrics
  - And prepare | wrangle metrics (tag, filter, enrich, aggregate, ...)
- To raise alert
  - To human (via ticket/SMS/...)
  - To automated handler/agent
- To support issue resolution (data for root cause analysis)
- To analyze trends + effects/impact of change

## Type of metrics



- Primary metrics
  - Tied to SLA indicators
  - Relevant – representative of Key Performance Indicators
    - End user experience
    - Business activity throughput
    - Cost efficiency
- Secondary metrics
  - Predictors for primary metrics
- Technical metrics
  - Hygiene factors – not directly tied to specific business indicators
  - Temperature, Remaining Free Storage, Network Load,

# Dashboard




- Provide starting point to respond to an alert – start investigation (drill down, query)
- Demonstration purposes
  - Always looks nice – a dashboard



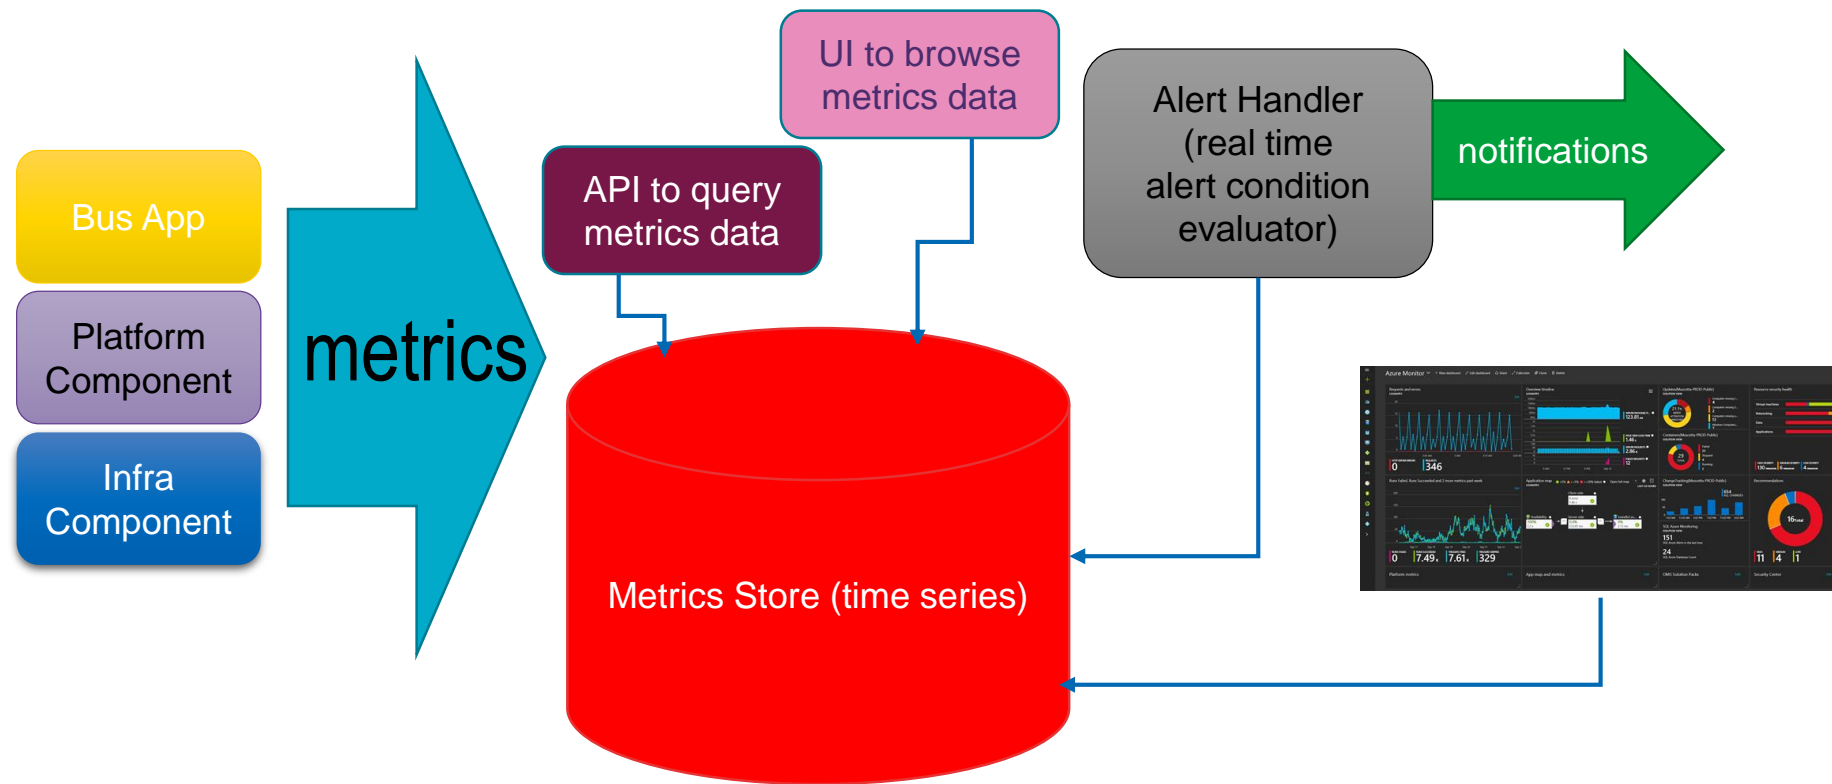
## Metrics are collected across the stack

- Business Applications
  - SaaS, Standard Applications
  - Custom | Tailor made applications
- Platform
  - Web Server, Application Server
  - Database
  - LDAP
  - JVM/Node/.NET/... runtime
- Infra
  - Container, Container Platform (e.g. Kubernetes)
  - Operating System
  - Cache
  - Proxy, Load Balancer
  - Network
  - Storage, File System





# Monitoring



# History and Status of Prometheus



- Written in Go Lang – all open source, available on GitHub
- Part of CNCF
- <https://prometheus.io/>
- Since 2012
- Defacto standard for gathering metrics (?)
- *treating time-series data as a data source for generating alerts is now accessible to everyone*



## Dimensional data

Prometheus implements a highly dimensional data model. Time series are identified by a metric name and a set of key-value pairs.



## Powerful queries

A flexible query language allows slicing and dicing of collected time series data in order to generate ad-hoc graphs, tables, and alerts.



## Great visualization

Prometheus has multiple modes for visualizing data: a built-in expression browser, Grafana integration, and a console template language.



## Efficient storage

Prometheus stores time series in memory and on local disk in an efficient custom format. Scaling is achieved by functional sharding and federation.



## Simple operation

Each server is independent for reliability, relying only on local storage. Written in Go, all binaries are statically linked and easy to deploy.



## Precise alerting

Alerts are defined based on Prometheus's flexible query language and maintain dimensional information. An alertmanager handles notifications and silencing.



## Many client libraries

Client libraries allow easy instrumentation of services. Over ten languages are supported already and custom libraries are easy to implement.



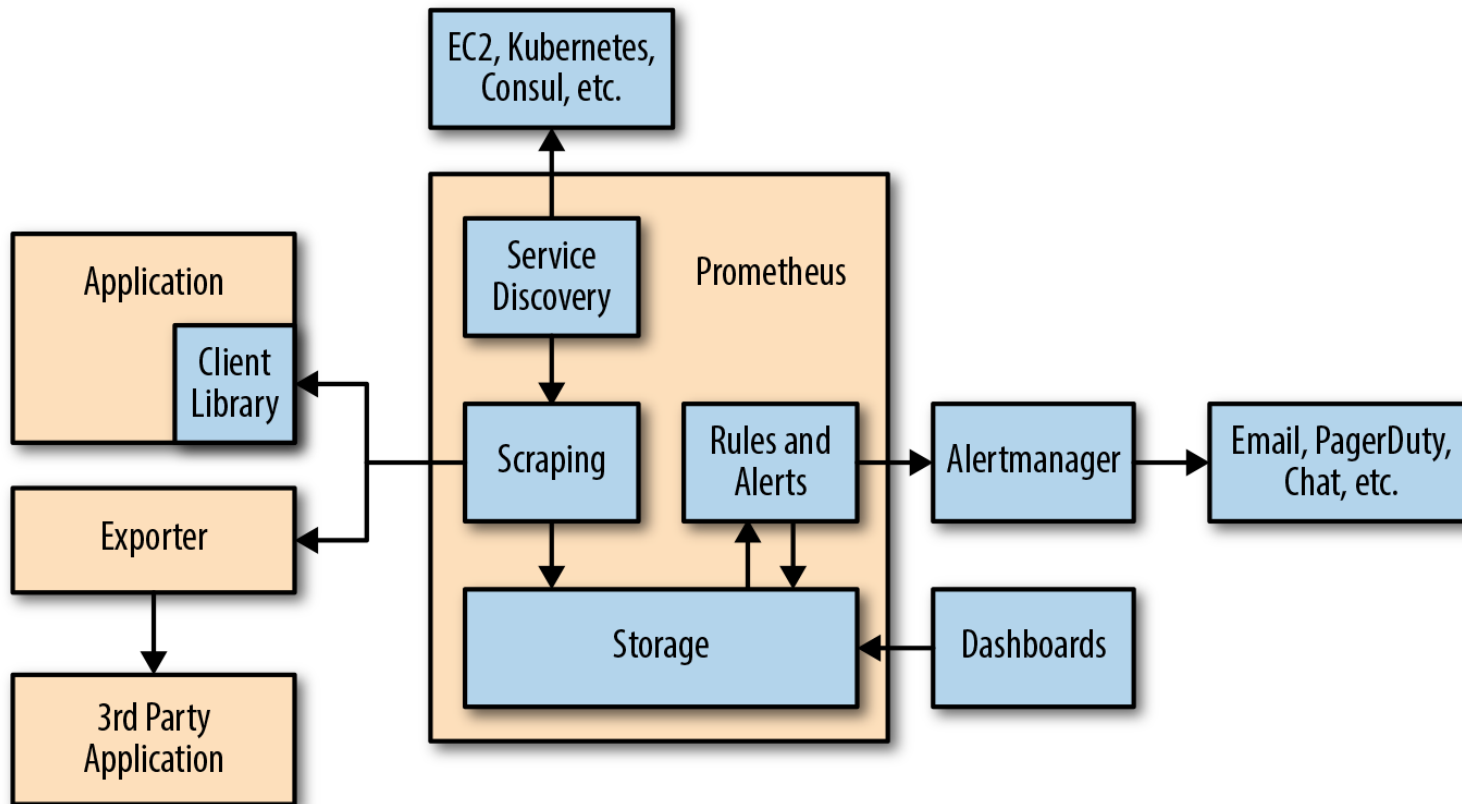
## Many integrations

Existing exporters allow bridging of third-party data into Prometheus. Examples: system statistics, as well as Docker, HAProxy, StatsD, and JMX metrics.

- Gather metrics into database
  - Scheduled pull [harvest] scrape actions – HTTP/TCP requests
  - Accessing Exporters and built in (scrape) endpoints
- Make metrics available to consuming systems and humans
  - Such as Grafana (for dashboarding)
  - REST APIs
  - Prometheus UI – Graphs, Console, PromQL
- Analyze metrics according to [alert] rules
  - Determine if alerts are “firing”
- Act on firing alerts
  - Send notifications
- Supports federation – global view over local environments and recovery of local environment



# Prometheus Architecture



- All are numeric
- Uniquely identified by name and set of labels
- All can be labeled (associated with dimensions)
  - Aggregation is done *grouped by dimension*
  - (labels should have limited number of values)
  - Filtering and Drill down is also done using labels
- Examples of labels:
  - Data Center, Region
  - Environment (Prod, Acc, Test)
  - Service, Application, Module
  - URL Path, request type, status, error code
  - Not for dynamic aspects such as date, time, user id, user IP

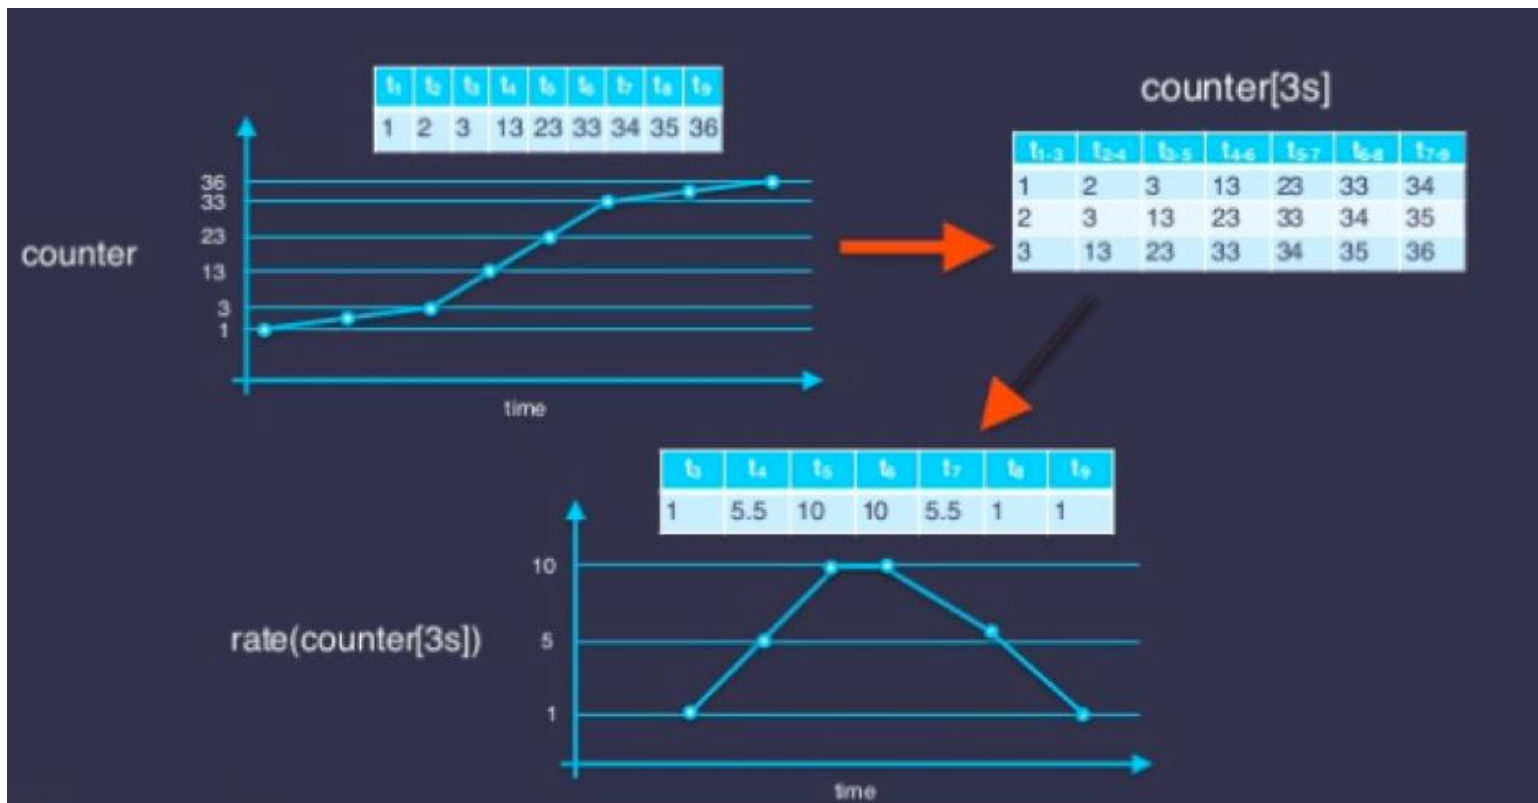
# Prometheus – Types of Metrics



- Basic counters:
  - Gauge – current value (%disk free, response time, state of database)
  - Counter - #occurrences – and the derived rate - #occurrences/time unit - (number of requests, number of bytes transferred);
    - counters can be reset but not decremented
- Sampling Counters:
  - Summary – reports quantiles as well as total sum and total count (over sliding window)
  - Histogram – maps observations to buckets and reports number of instances per bucket (over sliding window)


Basic Counters	Sampling Counters
counter	histogram
gauge	summary

## Counter – per time slice and as a rate



## Exposing Metrics for Prometheus to Scrape


- Have the application or environment listen for HTTP requests at a specific endpoint (for example: host:port/metrics)
- Return Metrics in the proper format to GET requests to this endpoint
- Use a Client Library to easily compose the proper metrics response messages
- Configure the endpoint on the Prometheus server in the *prometheus.yml* file





# Client Libraries for Exposing Metrics for Prometheus to Scrape from custom applications

- Go
- Java or Scala
- Python
- Ruby
- Bash
- C++
- Common Lisp
- Elixir
- Erlang
- Haskell
- Lua for Nginx
- Lua for Tarantool
- .NET / C#
- Node.js
- PHP
- Rust




## Demo – example of custom application with client library



# Prometheus Exporters

- Specialized adapters to expose metrics for specific technology components
  - Installed and configured for a specific component
  - Scraped by Prometheus based on a config file that references the exporter




# Prometheus Exporters

- Linux & Windows
  - Databases
  - Messaging Systems
  - Storage
  - HTTP
  - APIs
  - Logging
  - Monitoring Systems
  - Application Servers & Container Platforms
- 
- Built-in metric publication (without exporter)



## For example: Oracle Database Exporter



## Prometheus Oracle Exporter

A [Prometheus](#) exporter for Oracle.



The following metrics are exposed currently. Support for RAC (databasename and instancename added via labels)

- `oracledb_exporter_last_scrape_duration_seconds`
- `oracledb_exporter_last_scrape_error`
- `oracledb_exporter_scrapes_total`
- `oracledb_uptime (days)`
- `oracledb_session` (view `v$session` system/user active/passive)
- `oracledb_sysmetric` (view `v$sysmetric` (Physical Read Total IO Requests Per Sec / Physical Write Total IO Requests Per Sec / Physical Read Total Bytes Per Sec / Physical Write Total Bytes Per Sec))
- `oracledb_sysstat` (view `v$sysstat` (parse count (total) / execute count / user commits / user rollbacks))
- `oracledb_waitclass` (view `v$waitclass`)
- `oracledb_tablespace` (tablespace total/free)
- `oracledb_asmspace` (Space in ASM (`v$asm_disk/v$asm_diskgroup`))
- `oracledb_interconnect` (view `v$sysstat` (gc cr blocks served / gc cr blocks flushed / gc cr blocks received))
- `oracledb_recovery` (percentage usage in FRA from `V$RECOVERY_FILE_DEST`)
- `oracledb_redo` (Redo log switches over last 5 min from `v$log_history`)
- `oracledb_cachehitratio` (Cache hit ratios (`v$sysmetric`))
- `oracledb_up` (Whether the Oracle server is up)
- `oracledb_error` (Errors parsed from the alert.log)
- `oracledb_error_unix_seconds` (Last modified Date of alert.log in Unixtime)
- `oracledb_services` (Active Oracle Services (`v$active_services`))
- `oracledb_parameter` (Configuration Parameters (`v$parameter`))
- `oracledb_query` (Self defined Queries in Configuration File)

\*TOOK VERY LONG, BE CAREFUL (Put the Metrics below in a separate Scrape-Config):

- `oracledb_tablerows` (Number of Rows in Tables)
- `oracledb_tablebytes` (Bytes used by Table)
- `oracledb_indexbytes` (Bytes used by Indexes of associated Table)
- `oracledb_lobbytes` (Bytes used by Lobs of associated Table)

THE ORACLE EXPORTER FOR PROMETHEUS IS AVAILABLE AT <https://github.com/brimble/oracle-exporter>



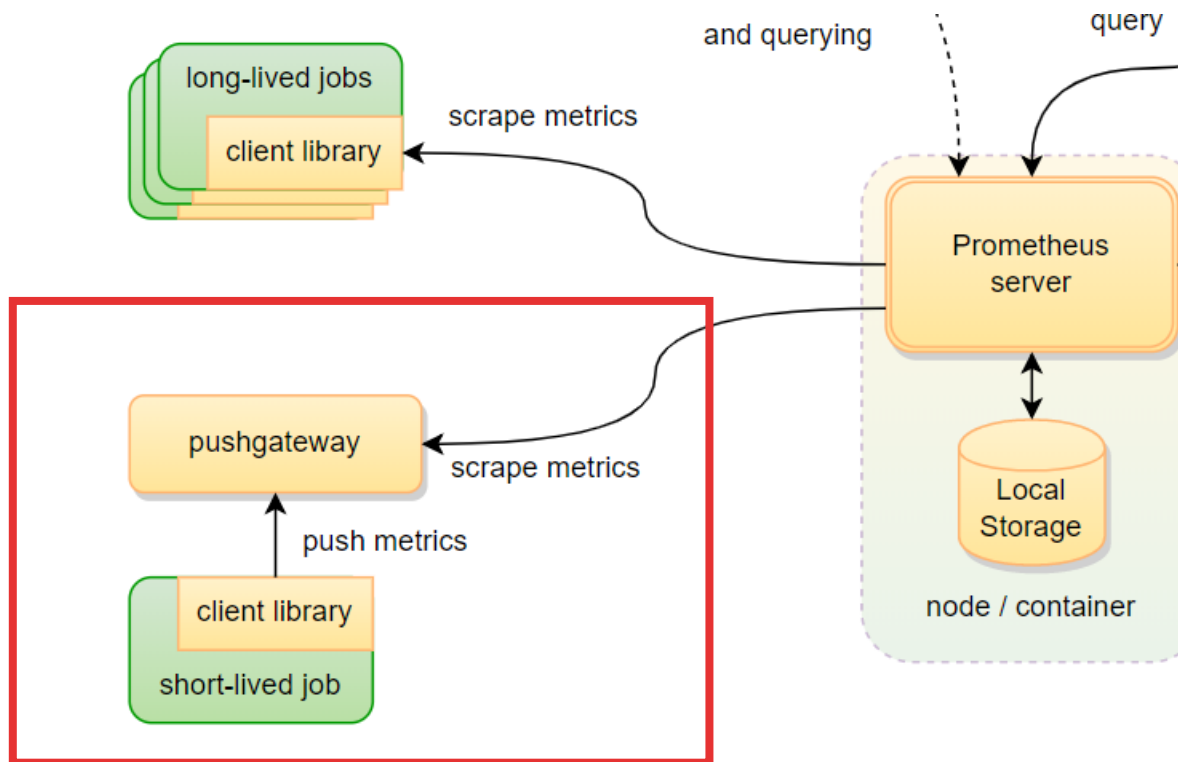
## Demo – (Linux) Node Exporter



- Run node exporter as background job on Linux node

## Pushgateway for Short-Lived Jobs

- Jobs that may be gone before their metrics are scraped





## Grafana - dashboarding




- Grafana is a generic dashboard product
  - Supporting many types of data sources, of which Prometheus is but one

## Getting Started/Useful Resources

- KataCoda:

<https://www.katacoda.com/courses/prometheus>

<https://www.katacoda.com/courses/prometheus/creating-dashboards-with-grafana>





- **Blog:** [technology.amis.nl](https://technology.amis.nl)
- **Email:** [lucas.jellema@amis.nl](mailto:lucas.jellema@amis.nl)
-  : [@lucasjellema](https://twitter.com/lucasjellema)
-  : [lucas-jellema](https://www.linkedin.com/in/lucas-jellema)
- **AMIS** : [www.amis.nl](https://www.amis.nl), [info@amis.nl](mailto:info@amis.nl)

धन्यवाद

dhanyavaad

Thank you

Dank je wel



<https://github.com/lucasjellema>